

EPGSpeller: Open-Vocabulary Silent Spelling with Electropalatography

Anonymous Authors

February 19, 2026

Abstract

We study open-vocabulary (lexicon-free, character-level) silent spelling from electropalatography (EPG) under a word-holdout protocol. We ask whether injecting a fixed proxy electrode layout as a spatial inductive bias improves character-level decoding. We compare (i) a layout-agnostic channel vector front-end, (ii) a proxy-grid projection consumed by a two-dimensional convolutional front-end, and (iii) row/column pooling on the same proxy grid. In this setting, the proxy-grid convolutional front-end has higher (worse) greedy CER than the vector baseline and a higher (worse) RTF, while proxy-grid spatial augmentation reduces CER degradation under fixed electrode dropouts.

1 Problem setting

Electropalatography (EPG) measures tongue–palate contact as a multichannel time series from a fixed electrode array. EPG has long been used to study and visualize tongue–palate contact patterns in phonetic research and speech training (e.g., [9, 10, 16]), and prior work has explored EPG data reduction and representation [11, 2]. We work with a fixed inventory of 124 electrode channels indexed by the SmartPalate channel numbering (SmartPalate refers to the EPG device and its channel indexing scheme), and we use that index order as the canonical channel order for layout-agnostic models.

To introduce spatial structure, we define a fixed proxy mapping π from SmartPalate channel indices to grid coordinates on a 16×16 grid, derived from the pinned SmartPalate distribution snapshot

`ev_scripts_smartpalate_distribution_csv` (`d21d93c9ec3c`). Operationally, the snapshot is a 16×16 table of channel indices; $\pi(c)$ is the first row/column location of channel index c in that table. Channels absent from the snapshot are treated as *unmapped* and retained as separate non-grid features (4); 5 channel indices appear multiple times in the snapshot and are resolved deterministically by first occurrence. We also provide a separate binary availability mask channel indicating which proxy-grid locations correspond to available channels under the chosen subset. This proxy mapping is used only as an inductive bias for layout-aware representations; we do not interpret it as a geometric tongue model.

We use a character-level connectionist temporal classification (CTC) formulation [8]. In the word-holdout protocol, training and test sets are disjoint at the word-identity level, where word identity is defined as the target-word string used for evaluation; all examples of each test word are excluded from training. We decode with greedy CTC and do not apply a lexicon constraint or an external language model.

2 Related work and positioning

Silent speech interfaces (SSI) decode linguistic content from non-acoustic signals; surveys summarize modality taxonomies and highlight protocol dependence across systems [3, 5, 7, 14]. SSI has been explored with multiple sensing modalities, including ultrasound and optical tongue imaging [12] and implanted magnetic-sensor setups [6]; here we focus on electropalatography (EPG), whose palatogram-style structure motivates explicit spa-

tial representations and reduction analyses [16, 11, 2, 17]. Closest to our task, SilentSpeller studies EPG-based silent spelling and text entry [13], while ReHEarSSE provides a recent non-EPG silent-spelling reference point [4]; EPG has also been used in multimodal learning beyond text entry [1]. In this work, under a fixed word-holdout protocol and greedy CTC decoding, we isolate the effect of proxy-layout-aware front-ends (proxy-grid two-dimensional convolution or row/column compression) and evaluate robustness to missing electrodes via structured masking and augmentation [15]. We therefore scope claims to this protocol and proxy-layout assumption, and do not claim new hardware, a new electrode geometry, or a new decoding objective.

3 Models and spatial front-ends

All variants share the same temporal backbone and CTC decoding setup. The temporal backbone produces 64-dimensional per-frame embeddings and uses the same recurrent stack across conditions (uni-gru; 5layers; 512hidden units). We apply a temporal unfolding operator (sliding-window chunking) with kernel length 32 and stride 4 before the recurrent stack; we vary only the EPG spatial front-end described below. For budgeted settings, we restrict the input to a subset of size K ; for each K , the same selected channels are provided to all compared spatial front-ends. In fixed-dropout evaluations, we additionally set a fixed subset of those available channels to zero for all samples in the run while leaving the availability mask unchanged.

Vector (layout-agnostic). The vector baseline uses the canonical channel index order and does not use any electrode layout information.

Layout-aware proxy-grid projection (two-dimensional convolution). We scatter the selected channels into the 16×16 proxy grid to form a two-channel tensor (contact value and a binary availability mask), and apply a two-dimensional convolutional front-end to obtain per-frame embeddings. Channels absent from the proxy-layout snapshot are treated as *unmapped* and embedded separately; duplicate indices in the snapshot are resolved by the definition of π (first occurrence) before scattering. We

combine the proxy-grid embedding and the unmapped embedding by summation after linear projection to the shared embedding dimension.

Row/column compression. We form row-wise and column-wise pooled summaries from the same proxy grid, concatenate them with the unmapped channels, and project to the shared embedding size.

4 Experimental protocol

We evaluate the word-holdout protocol using 4 random split seeds (seed0, seed1, seed2, seed3) and reuse the exact same splits across all compared spatial front-ends. Training uses 10000 optimization steps at 10.0 ms frame shift and enables SpecAugment-inspired masking [15]. Here, SpecAugment-inspired masking refers to time- and channel-wise masking applied to the input sequences for all models, whereas *spatial augmentation* refers to proxy-grid perturbations applied only to the proxy-grid convolutional front-end.

We compare two electrode-selection baselines for budgeted subsets of size K : *TopK*, which selects the K highest-ranked channels under an occlusion-based importance score (performance degradation induced by masking a channel), and *FPS-diverse TopK*, which adds a proxy-grid spatial-diversity constraint via farthest-point sampling from a high-scoring candidate pool. For a channel c , the occlusion score is the increase in greedy CER when channel c is set to zero at evaluation time while holding the trained model fixed; *TopK* selects the K channels with the largest occlusion scores (mean over splits). *FPS-diverse TopK* then selects a subset by greedy farthest-point sampling over proxy-grid coordinates using Euclidean distance on the proxy grid, from a candidate pool of twice the budget size; if mapped candidates are insufficient, remaining channels (including unmapped) are filled by importance order. We treat both as baselines and do not claim novelty for these heuristics.

We report character error rate (CER) from greedy CTC decoding and streaming real-time factor (RTF). CER is computed as total edit distance divided by total reference characters on the evaluation set. RTF is computed as total model forward-pass wall time (excluding decoding and data loading) divided by total input duration (derived from

Table 1: Vector vs two-dimensional convolution (mean $\pm$ std over splits). Evidence: `ev_ed20260217_h11_spatial2d_vs_vector_1` (2f2840f6dee5).

K	Vector Greedy CER	Spatial2D Greedy CER	Vector Stream RTF	Spatial2D Stream RTF
TopK				
32	0.1399 ± 0.0162	0.4885 ± 0.1269	0.000491 ± 0.000035	0.002670 ± 0.000127
64	0.1289 ± 0.0135	0.2609 ± 0.0300	0.000476 ± 0.000004	0.002895 ± 0.000304
96	0.1194 ± 0.0173	0.2304 ± 0.0350	0.000475 ± 0.000023	0.002767 ± 0.000238
124	0.1048 ± 0.0093	0.1831 ± 0.0103	0.000469 ± 0.000014	0.002639 ± 0.000067
FPS2K				
32	0.1440 ± 0.0136	0.5795 ± 0.3052	0.000514 ± 0.000022	0.002674 ± 0.000129
64	0.1195 ± 0.0103	0.3581 ± 0.3079	0.000532 ± 0.000094	0.002808 ± 0.000180
96	0.1137 ± 0.0161	0.3438 ± 0.3104	0.000486 ± 0.000029	0.002626 ± 0.000096
124	0.1048 ± 0.0093	0.1831 ± 0.0103	0.000469 ± 0.000014	0.002639 ± 0.000067

the number of input frames and the 10.0 ms frame shift). Lower is better for both metrics.

5 Results

Scope. All comparisons below use the same protocol and matched training settings; we interpret the findings as evidence about proxy-layout spatial inductive biases under this setting, not as claims about true geometric tongue–palate structure.

5.1 Two-dimensional convolution vs vector

Table 1 compares vector and two-dimensional convolutional front-ends for multiple electrode budgets. In this protocol and setting, the proxy-grid convolutional front-end has higher (worse) greedy CER than the vector baseline and a higher (worse) RTF.

5.2 Row/column compression vs vector

Table 2 shows that row/column compression yields similar greedy CER to the vector baseline, but a higher (worse) RTF.

5.3 Spatial augmentation and robustness to electrode failures

To quantify robustness, we apply fixed electrode-dropout masks at evaluation time (the same masked channels for all test samples in an evaluation run) and report the increase in greedy CER relative to clean evaluation. Here,

Table 2: Vector vs row/column compression (mean±std over splits). Evidence: `1e93ed2626269267k123_2026e02v37vector_p1_k(a7fb56ec370d)`.

K	Vector Greedy CER	RowCol Greedy CER	Vector Stream RTF	RowCol Stream RTF
TopK				
32	0.1399 ± 0.0162	0.1448 ± 0.0204	0.000491 ± 0.000035	0.002141 ± 0.000276
64	0.1289 ± 0.0135	0.1269 ± 0.0143	0.000476 ± 0.000004	0.002072 ± 0.000142
FPS2K				
32	0.1445 ± 0.0136	0.1466 ± 0.0080	0.000514 ± 0.000022	0.002176 ± 0.000311
64	0.1195 ± 0.0103	0.1255 ± 0.0193	0.000532 ± 0.000094	0.002104 ± 0.000216

Table 3: Fixed-dropout robustness for the two-dimensional front-end (mean $\pm$ std over splits). Values are Δ CER relative to clean evaluation. Evidence: `ev_ed20260217_h12_spatial_aug_fixed_dropout` (0cd1fecedc16).

Drop rate	No spatial aug	Spatial aug
0.0	0.0000 ± 0.0000	0.0000 ± 0.0000
0.1	0.0245 ± 0.0027	0.0183 ± 0.0065
0.2	0.0783 ± 0.0162	0.0660 ± 0.0097
0.3	0.1440 ± 0.0183	0.1224 ± 0.0120

spatial augmentation refers to training-time masking or perturbations that operate on the proxy electrode grid (rather than treating channels as exchangeable). Table 3 shows that spatial augmentation reduces the CER degradation under dropout.

6 Artifacts and auditability

This project maintains a claim-evidence registry in `paper/paper.json`. All numeric facts and result tables in this draft are generated from pinned evidence snapshots via `scripts/paper/generate_facts_tex.py` using `paper/facts_spec.yaml`. A strict checker rejects handwritten numeric facts in the manuscript body. To regenerate the manuscript numbers from a clean checkout, run: `python3scripts/paper/generate_facts_tex.py--facts_specpaper/facts_spec.yaml--paper_jsonpaper/paper.json--root.` and `python3scripts/paper/check_no_manual_`

```
numeric_facts.py--texmain.tex--facts_
specpaper/facts_spec.yaml--root.. To au-
dit claim-evidence consistency, run: python3tools/
paper_json_lint.py--paper-jsonpaper/
paper.json--root.--strict and
python3scripts/render_paper_json_
claim_audit.py--paper-jsonpaper/paper.
json--root.--out-mdpaper/snippets/
claim_audit.md--strict.
```

Key fiers	pinned used	evidence by	identi- draft	include
ev_scripts_smartpalate_distribution_csv	(d21d93c9ec3c)	ev_ed20260217_h11_spatial2d	vs_vector_p1_k32_k64_k96_k124_2026-02-17	
(2f2840f6dee5)	ev_ed20260217_h13_rowcol	vs_vector_p1_k32_k64_2026-02-17		
(a7fb56ec370d)				and
ev_ed20260217_h12_spatial_aug_fixed_dropout	(0cd1fecedc16)		summary_p1_topk_k64_2026-02-17	

References

- [1] arxiv: Epg2s: Speech generation and speech enhancement based on electropalatography and audio signals using multimodal learning, 2022.
- [2] Miguel Á. Carreira-Perpiñán and Steve Renals. Dimensionality reduction of electropalatographic data using latent variable models, 1998.
- [3] B. Denby, T. Schultz, K. Honda, T. Hueber, J.M. Gilbert, and J.S. Brumberg. Silent speech interfaces, 2010.
- [4] Xuefu Dong, Yifei Chen, Yuuki Nishiyama, Kaoru Sezaki, Yuntao Wang, Ken Christofferson, and Alex Mariakakis. Rehearsse: Recognizing hidden-in-the-ear silently spelled expressions, 2024.
- [5] João Freitas, António Teixeira, Miguel Sales Dias, and Samuel Silva. *An Introduction to Silent Speech Interfaces*. 2017.
- [6] J.M. Gilbert, S.I. Rybchenko, R. Hofe, S.R. Ell, M.J. Fagan, R.K. Moore, and P. Green. Isolated word recognition of silent speech using magnetic implants and sensors, 2010.
- [7] Jose A. Gonzalez-Lopez, Alejandro Gomez-Alanis, Juan M. Martin Donas, Jose L. Perez-Cordoba, and Angel M. Gomez. Silent speech interfaces for speech restoration: A review, 2020.
- [8] Alex Graves, Santiago Fernández, Faustino Gomez, and Jürgen Schmidhuber. Connectionist temporal classification, 2006.
- [9] W. Hardcastle, W. Jones, C. Knight, A. Trudgeon, and G. Calder. New developments in electropalatography: A state-of-the-art report, 1989.
- [10] William J. Hardcastle. Electropalatography in phonetic research and in speech training, 1990.
- [11] W.J. Hardcastle, F. Gibbon, and K. Nicolaidis. Epg data reduction methods and their implications for studies of lingual coarticulation, 1991.
- [12] Thomas Hueber, Elie-Laurent Benaroya, Gérard Chollet, Bruce Denby, Gérard Dreyfus, and Maureen Stone. Development of a silent speech interface driven by ultrasound and optical images of the tongue and lips, 2010.
- [13] Naoki Kimura, Tan Gemicioglu, Jonathan Womack, Richard Li, Yuhui Zhao, Abdelkareem Bedri, Zixiong Su, Alex Olwal, Jun Rekimoto, and Thad Starner. Silentspeller: Towards mobile, hands-free, silent speech text entry using electropalatography, 2022.
- [14] Wookey Lee, Jessica Jiwon Seong, Busra Ozlu, Bong Sup Shim, Azizbek Marakhimov, and Suan Lee. Biosignal sensors and deep learning-based speech recognition: A review, 2021.
- [15] Daniel S. Park, William Chan, Yu Zhang, Chung-Cheng Chiu, Barret Zoph, Ekin D. Cubuk, and Quoc V. Le. SpecAugment: A simple data augmentation method for automatic speech recognition, 2019.
- [16] Jo Verhoeven, Naomi Rachel Miller, Luc Daems, and Constantino Carlos Reyes-Aldasoro. Visualisation and analysis of speech production with electropalatography, 2019.

- [17] Seong-Tak Woo, Ji-Wan Ha, Sungdae Na, Hyunjoon Choi, and Sung-Bom Pyun. Design and evaluation of korean electropalatography (k-epg), 2021.